

# 1 Changes introduced in HDF5\_Parallel Branch

Items listed in blue effect compatibility of older code when using 3.2.2. Known problems are highlighted in red.

## 1.1 General behavior changes and new recommendations for parallel performance

- The flush functions should not be used. Staggering the writing and reading by writing the data immediately avoids IO contention occurring when flush is being used.
- The parallel routines are meant for parallel file systems (GPFS or Lustre).
- The default parallel input/output mode was changed from CGP\_INDEPENDENT to CGP\_COLLECTIVE.
- An extra argument for passing MPI info to the CGNS library was added to `cgp_pio_mode`.

### C

```
int cgp_pio_mode(CGNS_ENUMT(PIOmode_t) mode, MPI_Info info)
```

### Fortran

```
CALL cgp_pio_mode_f(mode, comm_info, ierr)
INTEGER(KIND(CGP_COLLECTIVE)) :: mode ! Use parameters CGP_INDEPENDENT or CGP_COLLECTIVE
INTEGER :: comm_info
INTEGER :: ierr
```

- Functions for parallel reading and writing multi-component datasets using a single call was introduced. The new APIs use new capabilities introduced in version 1.8.\* (currently not released as of 10.13.2014) of the HDF5 library. The new APIs pack multiple datasets into a single buffer and the underlining MPI IO completes the IO request using just one call. The availability of the new functions in the HDF5 library is checked at compile time. The current limitation (due to MPI) is that the size of the datasets must be less than 2GB.

### C

```
int cgp_coord_multi_read_data(int fn, int B, int Z, int *C,
                             const cgsz_t *rmin, const cgsz_t *rmax,
                             void *coordsX, void *coordsY, void *coordsZ);

int cgp_coord_multi_write_data(int fn, int B, int Z, int *C,
                              const cgsz_t *rmin, const cgsz_t *rmax,
                              const void *coordsX, const void *coordsY, const void *coordsZ);

int cgp_field_multi_read_data(int fn, int B, int Z, int S, int *F,
                             const cgsz_t *rmin, const cgsz_t *rmax,
                             int nsets, ...);
/* ... nsets of variable arguments, *solution_array, corresponding to the order given by F */

int cgp_field_multi_write_data(int fn, int B, int Z, int S, int *F,
                              const cgsz_t *rmin, const cgsz_t *rmax,
                              int nsets, ...);
/* ... nsets of variable arguments, *solution_array, corresponding to the order given by F */

int cgp_array_multi_write_data(int fn, int *A, const cgsz_t *rmin, const cgsz_t *rmax,
                              int nsets, ...);
/* ... nsets of variable arguments, *field_array, corresponding to the order given by F */

int cgp_array_multi_read_data(int fn, int *A, const cgsz_t *rmin, const cgsz_t *rmax,
                              int nsets, ...);
/* ... nsets of variable arguments, *field_array, corresponding to the order given by F */
```

## Fortran

```
CALL cgp_coord_multi_read_data_f(fn, B, Z, C, rmin, rmax, coordsX, coordsY, coordsZ, ier)
  INTEGER :: fn
  INTEGER :: B
  INTEGER :: Z
  INTEGER :: C
  INTEGER(CG_SIZE_T) :: rmin
  INTEGER(CG_SIZE_T) :: rmax
  REAL :: coordsX, coordsY, coordsZ
  INTEGER :: ier
```

```
CALL cgp_coord_multi_write_data_f(fn, B, Z, C, rmin, rmax, coordsX, coordsY, coordsZ, ier)
  INTEGER :: fn
  INTEGER :: B
  INTEGER :: Z
  INTEGER :: C
  INTEGER(CG_SIZE_T) :: rmin
  INTEGER(CG_SIZE_T) :: rmax
  REAL :: coordsX, coordsY, coordsZ
  INTEGER :: ier
```

```
CALL cgp_field_multi_write_data_f(fn, B, Z, S, F, rmin, rmax, ier, nsets, ...)
  INTEGER :: fn
  INTEGER :: B
  INTEGER :: Z
  INTEGER :: C
  INTEGER(CG_SIZE_T) :: rmin
  INTEGER(CG_SIZE_T) :: rmax
  INTEGER :: ier
  INTEGER :: nsets
  ... REAL, DIMENSION(*) :: field_array ! entered nsets times
```

```
CALL cgp_field_multi_read_data_f(fn, B, Z, S, F, rmin, rmax, ier, nsets, ...)
  INTEGER :: fn
  INTEGER :: B
  INTEGER :: Z
  INTEGER :: C
  INTEGER(CG_SIZE_T) :: rmin
  INTEGER(CG_SIZE_T) :: rmax
  INTEGER :: ier
  INTEGER :: nsets
  ... REAL, DIMENSION(*) :: field_array ! entered nsets times
```

```
CALL cgp_array_multi_write_data_f(fn, B, Z, S, F, rmin, rmax, ier, nsets, ...)
  INTEGER :: fn
  INTEGER :: B
  INTEGER :: Z
  INTEGER :: C
  INTEGER(CG_SIZE_T) :: rmin
  INTEGER(CG_SIZE_T) :: rmax
  INTEGER :: ier
  INTEGER :: nsets
  ... REAL, DIMENSION(*) :: data_array ! entered nsets times
```

```
CALL cgp_array_multi_read_data_f(fn, B, Z, S, F, rmin, rmax, ier, nsets, ...)
INTEGER :: fn
INTEGER :: B
INTEGER :: Z
INTEGER :: C
INTEGER(CG_SIZE_T) :: rmin
INTEGER(CG_SIZE_T) :: rmax
INTEGER :: ier
INTEGER :: nsets
... REAL, DIMENSION(*) :: data_array ! entered nsets times
```

## 1.2 New C changes

- A new example benchmark program, benchmark\_hdf5.c was added to ptests.

## 1.3 New Fortran changes

All users are **strongly** encouraged to use a Fortran 2003 standard compliant compiler. Using a Fortran 2003 compiler guarantees interoperability with the C APIs via the ISO\_C\_BINDING module. Many changes were added to the CGNS library in order to take full advantage of the interoperability offered by the ISO\_C\_BINDING module.

1. Configure was changed to check if the Fortran compiler is Fortran 2003 compliant. If it is then the features of ISO\_C\_BINDING will be used.
2. The predefined CGNS constant parameters data types were changed from INTEGER to ENUM, BIND(C) for better C interoperability. The users should use the predefined types whenever possible and not the numerical value represented by the constant.
3. *INCLUDE "cgslib\_h"* was changed to using a module, USE CGNS.
  - (a) This allows defining a KIND type for the integers instead of the current way of using the preprocessor dependent *cgsizet*.
  - (b) Compatibility might be added back before the merge to the trunk.
4. The user should be sure to declare the arguments declared *int* in the C APIs as INTEGER in Fortran. The ONLY fortran arguments declared as type *cgsizet* should be the arguments also declared *cgsizet* in the C APIs. This is very important when building with option *-enable-64bit*.
5. Assuming the rules in step 4 were followed, users should not need to, or use, the CG\_BUILD\_64BIT parameters since Fortran's *cgsizet* is now guaranteed to match C's *cgsizet*.
6. Fortran programs defining CGNS data types with a default INTEGER size of 8 bytes are not currently compatible. This is independent of whether or not *-enable-64bit* is being used. For clarification, using *-enable-64bit* allows for data types (i.e. those declared as *cgsizet*) to be able to store values which are too large to be stored as 4 byte integers (i.e. numbers greater than 2,147,483,647). It is not necessary, or advisable, to have CGNS INTEGER types (types declared *int* in C) to be 8 bytes; the variables declared as *cgsizet* will automatically handle data types that can not be stored as 4 byte integers when *-enable-64bit* is being used.
  - (a) CGNS developer's note: A new C data type, *cgint\_f*, was introduced to be interpretable with the C type *int*. In order to allow for default 8 byte integers in Fortran: (1) The C API wrappers in *cg\_ftoc.c* need to be changed from *cgsizet* to *cgint\_f* everywhere the C argument is declared as an *int* in C, (2) configure needs to detect what size the default integer is in Fortran and find the corresponding size in C in order to set the correct size of *cgint\_f*.
7. Two new benchmarking programs were introduced in directory ptests:
  - (a) benchmarking\_hdf5\_f90.F90 uses the conventional Fortran wrappers.
  - (b) benchmarking\_hdf5\_f03.F90 calls the C APIs directly, not Fortran wrappers are used.

## 2 Parallel installation instructions

Two parallel files systems have been investigated: GPFS (mira, Argonne National Laboratory) and Lustre (Pleiades NASA). The following descriptions were performed on these systems, but they procedure should be similar for different machines of the same type. Example build scripts for these systems can be found in `src/sampleScripts`. They include scripts for building `zlib`, `hdf5` (assuming the user does not already have them install system wide) and a script for building `CGNS`. All the scripts use `autotools`, `cmake` remains untested. The next few examples assumes all the needed packages are in `${HOME}/packages` and all the scripts are placed in `${HOME}/packages`. This information can also be found in the `README.txt` in the scripts directory.

### 2.1 Building on IBM Blue Gene (GPFS)

1. Building `zlib` from source: Download and extract the `zlib` source: <http://www.zlib.net/>
  - (a) `cd` into the top level `zlib` source directory.
  - (b) modify and run the script: `../build_zlib`
2. Building `hdf5` from source
  - (a) From the top level of the `hdf5` library, change the `${HOME}/packages` to where `zlib` was installed in STEP 1.
  - (b) `../build_hdf5 -without-pthread -disable-shared -enable-parallel -enable-production \ -enable-fortran -enable-fortran2003 \ -disable-stream-vfd -disable-direct-vfd \ -with-zlib=${HOME}/packages/zlib-1.2.8/lib -prefix=${HOME}/packages/phdf5-trunk`

where `prefix` is set for where the `hdf5` library will get installed. There should be no need to modify in the script.

3. Building `cgns` from source:
  - (a) `cd` into the `cgns/src` directory
  - (b) modify and run: `<pathto>/build_cgns`
  - (c) `make`
  - (d) To make the tests: `cd ptests; make; make tests`
4. IMPORTANT PARAMETERS FOR GOOD PERFORMANCE
  - (a) The environment variable `BGLOCKLESSMPIO_F_TYPE=0x47504653` should be set. For example, this can be set using `qsub -env BGLOCKLESSMPIO_F_TYPE=0x47504653`

### 2.2 Building on SGI (Lustre)

## 3 General installation instructions

1. Install `HDF5` on your system
  - (a) `HDF5` uses the standard GNU `autotools`, so `'./configure'`, `'make'`, `'sudo make install'` should install `HDF5` without problems on most systems.
2. Unpack the tar ball containing the source code into some directory.
3. Create a new director in which to build the library.
4. Use `cmake` to initialize the build tree.

```
user@hostname:build_path$ cmake /path/to/cgns/sources/
```

5. Use cmake to edit the control variables as needed.

```
user@hostname:build_path$ cmake .
```

- (a) The paths to the HDF5 libraries and header files must be set in 'HDF5\_LIBRARY\_DIR' and 'HDF5\_INCLUDE\_' respectively.
  - i. If HDF5 is built with parallel-IO support via MPI, the 'HDF5\_NEEDS\_MPI' flag must be set to true.
  - ii. If HDF5 is built with zlib and szip support, these need to be flagged with 'HDF5\_NEEDS\_ZLIB' and 'HDF5\_NEEDS\_SZIP' as well as the paths for those libraries.
- (b) Fortran can be enabled by toggling the 'CGNS\_ENABLE\_FORTRAN' variable.
  - i. A view of the attempt to autodetect the correct interface between Fortran and C is show, setting the value of 'FORTRAN\_NAMING'.
  - ii. For gfortran and pgf90 the value of 'FORTRAN\_NAMING' should be 'LOWERCASE\_'.
- (c) The build system must be reconfigured after variable changes by pressing 'c'. Variables who's value has changed are marked with a '\*' in the interface.
- (d) After configuration, the Makefiles must be generated by pressing 'g'.

6. Use make to build the library.

```
user@hostname:build_path$ make
```

- (a) A colorized review of the build process should follow.

7. Installation of the library is accomplished with the 'install' target of the makefile.

```
user@hostname:build_path$ make install
```

- (a) You must have permissions to alter the directory where cgns will be installed.